



Data Management using mongoDB® & Mongoose



Course Roadmap



Web Client

Request

Response



Web Server

Frontend development

HTML for page structure



CSS for styling



JavaScript for interaction



Backend development

Web API



Data Management



Outline

1. [Introduction to MongoDB](#)
2. [Document Schema Design](#)
3. [Introduction to Mongoose](#)
4. [CRUD Operations](#)
5. [Aggregation Queries](#)

Introduction to

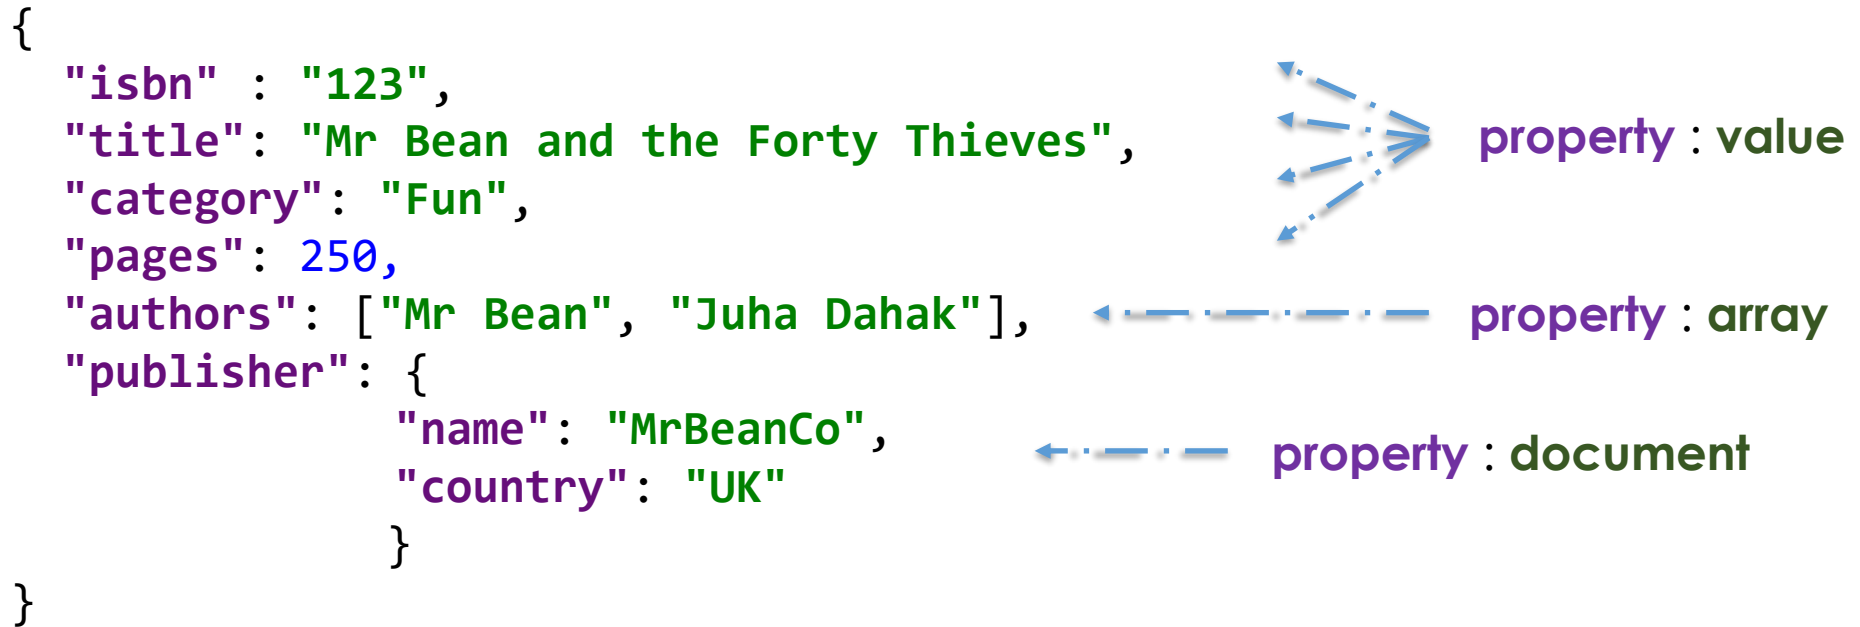


mongoDB®

What is MongoDB?

- MongoDB is an open-source **Document Oriented Database**
 - **Uses a document data model**: Stores data as JSON documents (instead of rows and columns as done in a relational database)
 - **Arrange documents in collections** (documents can vary in structure)
 - **API to query and manage documents**
- Better alternative data management solution for Web applications compared to using a Relational Database

Document



- **Document = JSON object**
- **Document = set of key-value pairs**
- **Basic unit of data** in MongoDB
- Analogous to **row** in a relational database

Collection

```
{
  "isbn": "123",
  "title": "Mr Bean and the Forty Thieves",
  "authors": ["Mr Bean", "Juha Dahak"],
  "publisher": {"name": "MrBeanCo", "country": "UK"},
  "category": "Fun",
  "pages": 250
}
```

- **Collection = Group** of documents
- Analogous to **table** in a relational database
- **Does not enforce** a schema
- Documents in a collection usually **have similar purpose** but they may have slightly different schema

Introduction to Mongoose




Mongoose Overview

- Mongoose is a Node.js **Object Document Mapper (ODM)** for MongoDB
 - Allows define **schemas to model** documents. Then use the **model** to read/write documents
 - A **schema** describes a document structure in terms of properties and their types.
 - You can add [validation](#), [virtual properties](#)
 - You can establish [references](#) to other models
 - A **model** is created based on a schema
 - A **model** maps to a MongoDB collection
 - A **model** = class used to run queries against collections
 - Instances of a model represent documents in MongoDB
 - Supports data validation on save
 - Allow rich querying of documents

MongoDB & Mongoose

MongoDB

Mongoose



Programming Steps

1. Import mongoose **module** `import mongoose from 'mongoose';`
2. Define a **schema** for each document = **Structure doc**

```
const storeSchema = new mongoose.Schema({  
  name: String,  
  city: String  
})
```

3. Create a **model** object based

```
const Store = mongoose.model('Store', storeSchema);
```

4. **Connect** to MongoDB

```
const dbConnection = mongoose.connect('mongodb://localhost/dbName');
```

5. Use the model to **read/write** documents

```
Store.find({})    //get all stores
```

Document Instance vs. Schema

```
{  
  "firstname" : "Simon",  
  "surname" : "Holmes",  
  "_id" : ObjectId("52279effc62ca8b0c1000007")  
}
```

**Example MongoDB
document**

```
{  
  firstname : String,  
  surname : String  
}
```

**Corresponding
Mongoose schema**

Schema Data Types

Example

Each property must have a type:

- String
- Number
- Date
- Boolean
- ObjectId
- Array

```
const reviewSchema = new mongoose.Schema({
  author: String,
  rating: {type: Number, required: true, min: 0, max: 5},
  reviewText: String,
  createdAt: {type: Date, default : Date.now}
})

const bookSchema = new mongoose.Schema({
  isbn: String,
  title: String,
  authors: [String],
  publisher: {name: String, country: String},
  category: String,
  pages: Number,
  read: {type: Boolean, default:false, required: true},
  createdAt : {
    type : Date,
    default : Date.now
  },
  reviews: [reviewSchema],
  store : [{ type : mongoose.Schema.ObjectId, ref : 'Store' }]
})
```

_id

- **_id** is the primary key that uniquely identifies each document in the collection
- It is automatically added when adding a document to the collection
- It is immutable (it cannot be changed)
- It is guaranteed to be unique across the whole database

Property Validation

- Built-in validators: required, min, max
- Can define custom validators

```
bookSchema.path('isbn').validate( value => value.length >= 3 )
```

- Validation happens on save

Virtual Property

- Define a property that won't get persisted to MongoDB

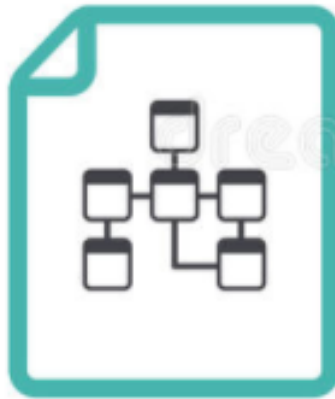
```
//define a fullName property that won't get persisted to MongoDB.  
personSchema.virtual('fullName').get(function () {  
    return `${this.name.first} ${this.name.last}`;  
});
```

```
// create a document  
const student1 = new Person({  
    name: { first: 'Ali', last: 'Faleh' }  
});
```

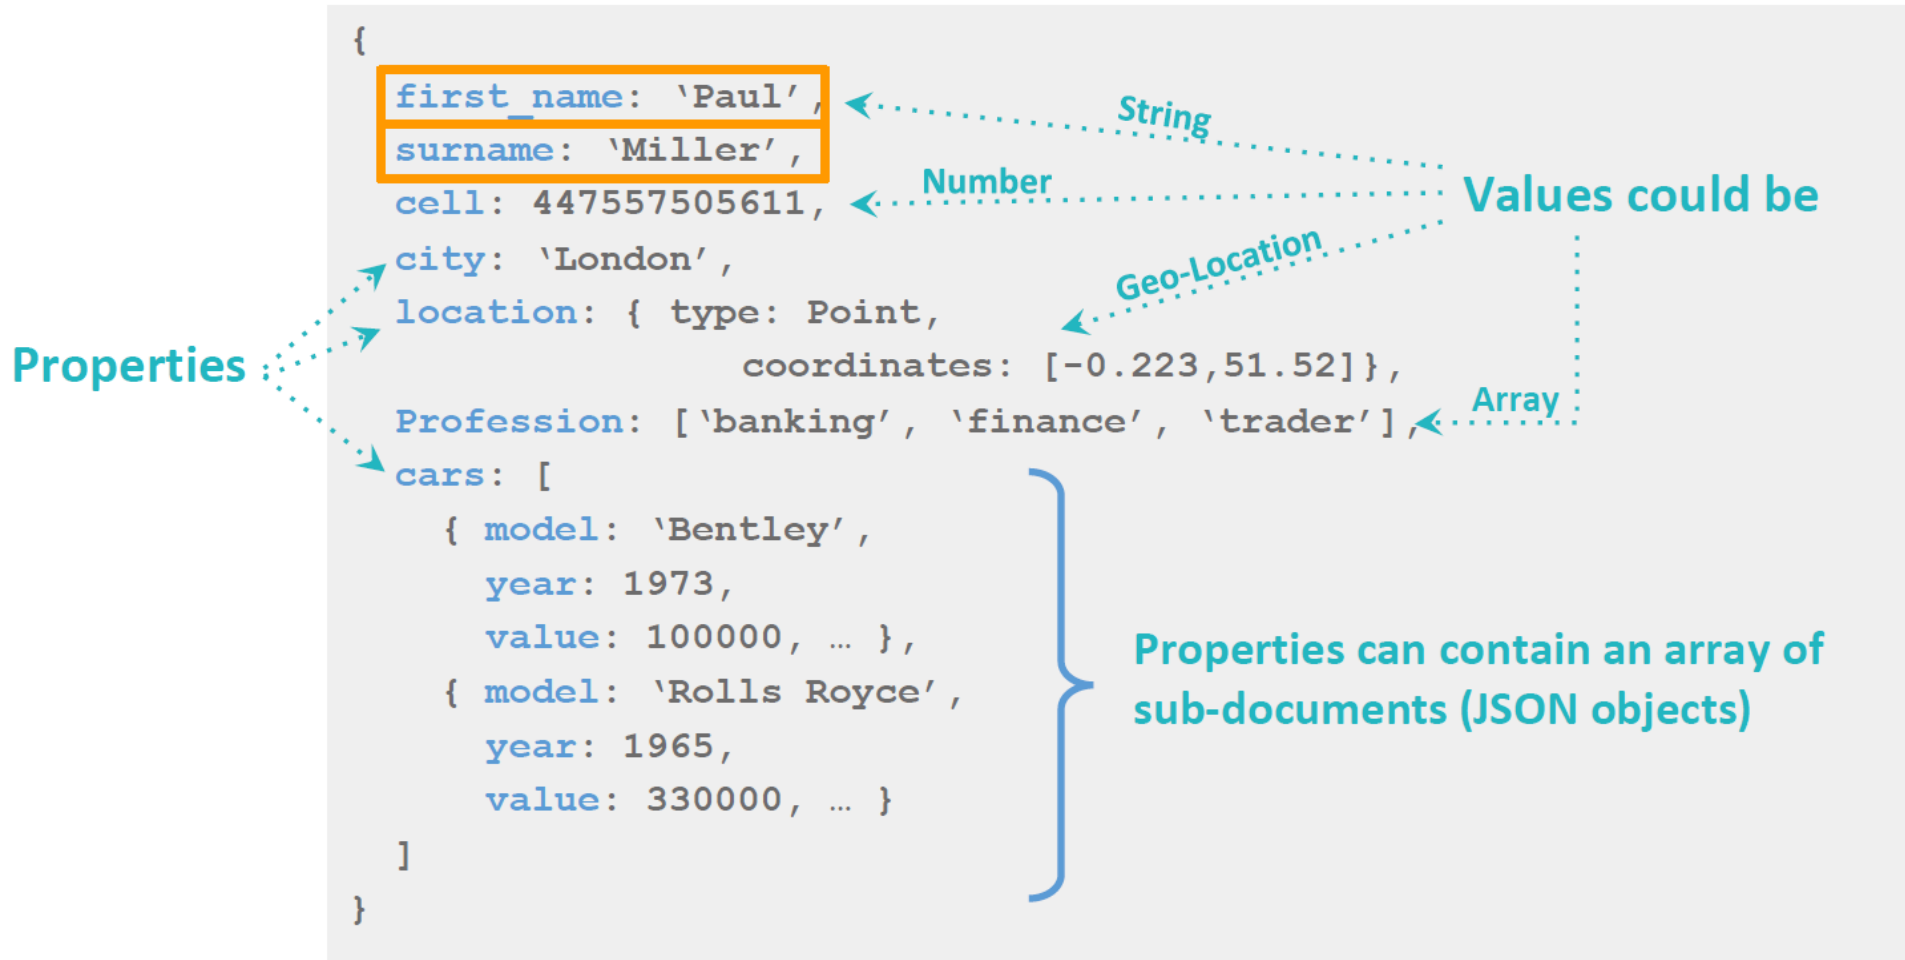
```
console.log(student1.fullName); // Ali Faleh
```

```
/*  
If you use toJSON() mongoose will not include virtuals by default  
unless if you pass { virtuals: true }  
*/  
console.log(student1.toJSON({ virtuals: true }));
```


Document Schema Design



Document can have a Complex Structure



Embedded vs Referenced documents

- Major design decision when designing a Document Schema is to decide **Embedded** vs **Referenced** subdocuments
- Decision should consider:
 - How the data will be used
 - Size of the document

One to Many Relationships

Patients

Embed

```
{
  _id: 2,
  first: "Joe",
  last: "Patient",
  addr: { ...},
  procedures: [
    {
      id: 12345,
      date: 2015-02-15,
      type: "CAT scan",
      ...},
    {
      id: 12346,
      date: 2015-02-15,
      type: "blood test",
      ...}]
}
```

OR

Patients

Reference

```
{
  _id: 2,
  first: "Joe",
  last: "Patient",
  addr: { ...},
  procedures: [12345, 12346]
}
```

Procedures

```
{
  _id: 12345,
  date: 2015-02-15,
  type: "CAT scan",
  ...}
{
  _id: 12346,
  date: 2015-02-15,
  type: "blood test",
  ...}
```

Embedding

- **Advantages**

- Retrieve all relevant information in a single query/document
- Avoid implementing joins in application code => fast data retrieval
- Update related information as a single atomic operation

- **Limitations**

- Large documents mean more overhead if most fields are not relevant
- 16 MB document size limit

Referencing

- **Advantages**


- Smaller documents
- Less likely to reach 16 MB document limit
- Infrequently accessed information not retruned on every query
- No duplication of data

- **Limitations**

- Two queries required to retrieve information

1 to 1 Relationships => Better to Embed

Medical Procedures

```
{  
   "_id": 333,  
  "date": "2003-02-09T05:00:00",  
  "hospital": "County Hills",  
  "patient": "John Doe",  
  "physician": "Stephen Smith",  
  "procedure": "Glucose",  
  "result": {  
    "value": 97,  
    "measurement": "mg/dl"  
  }  
}
```

Embed:

- No data duplication
- Data that are read/written together lives together

← Embed – *weak entity*

One to Many : General Recommendations

- **Embed when:**
 - **One-to-few** (e.g. customer - addresses)
 - **Often queried/updated together** in a single query (e.g., book - chapters)
 - No need to access the embedded object outside the context of the parent object (e.g., order – order items)
 - No additional data duplication introduced
- **Reference when:**
 - **1 to a large number of related items** (e.g. customer orders , book – reviews, video - comments)
 - Related data changes frequently (e.g., video – viewCount)
 - Referenced entity that is used by many others (e.g., session - **room**)
 - Document size is > 16 MB
 - Subdocument has a large number of infrequently accessed fields

1 to M Example 1

- *"We need to store user information like name, email and their addresses... a user they can have more than one address."*

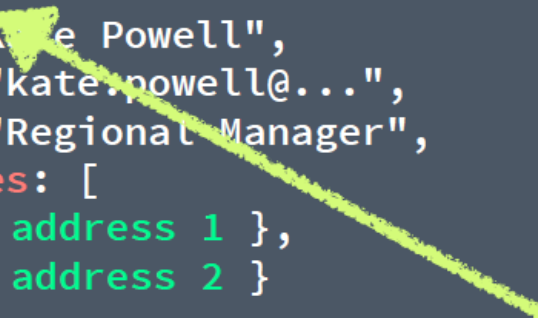
```
{
  _id: 1,
  name: "Kate Powell",
  email: "kate.powell@somedomain.com",
  title: "Regional Manager",
  addresses: [
    { street: "123 Sesame St", city: "Boston" },
    { street: "123 Evergreen St", city: "New York" }
  ]
}
```

One-to-few : embedding is the best design

1 to M Example 2

- "We have to be able to store tasks, assign them to users and track their progress..."*

```
> db.user.findOne({_id: 1})    > db.task.findOne({user_id: 1})
{                                {
  _id: 1,                        _id: 5,
  name: "Kate Powell",           summary: "Contact sellers",
  email: "kate.powell@...",      description: "Contact agents
  title: "Regional Manager",      to specify our ...",
  addresses: [                   due_date: ISODate(),
    { // address 1 },             status: "NOT_STARTED",
    { // address 2 }             user_id: 1
  ]
}
```



Referencing is the best design:

- **Tasks** are unbounded items: initially we do not know how many tasks we are going to have
 - A user can end with thousands of tasks
 - Maximum document size in MongoDB: 16 MB !
- Tasks can be queried without needing to retrieve the user details

CRUD Operations



CREATE



READ



UPDATE



DELETE

C

R

U

D

CRUD operations

- Create → `Book.create(newBook)`
- Read → `Book.find({})`
`Book.findById(bookId)`
`Book.findOne({isbn: "123"})`
`Book.find({authors: {$in: [author]}})`
- Update → `Book.update({_id: bookId}, updatedBook)`
- Delete → `Book.remove({_id : bookId});`

Mongoose Queries

- Queries are based on finding documents with any combination of fields in a collection

```
Book.find({ category: 'Fun', pages : { $lt : 200 } })
```

- You sorting and limits the number of returned documents

```
Book.find({}).sort('price').limit( 5 )
```

- OR condition is also supported

```
Book.find({}).where({ category: 'Fun' }).or({pages :{ $lt : 100 } })
```

- Filter on the existence of field

```
Book.find( { reviews : { $exists: true } } )
```

QueryBuilder

- The query object allows chaining methods could chained to build a complex query

```
School.find({ name: 'Iqraa'})  
.where('state').equals('AZ')  
.where('licenses').gt(17).lt(100)  
.where('district').in(['dist1', 'dist2'])  
.limit(10)  
.populate ('owner', 'name')  
.sort('owner.name')  
.select('id name state owner.name')
```

Count and Distinct Methods

- `collection.count(query)` - returns the number of documents in the collection that match the query
- `collection.distinct(field, query)` - returns an array of all the unique values found in the passed field for the documents that match the query


Populating a Reference Property

- Populating a reference property is the process of automatically **replacing** the reference Id is specified path in the document with document(s) from other collection(s)
- Populate sends another query for the related object

```
const bookSchema = new mongoose.Schema({  
  isbn: String,  
  title: String,  
  ...  
  stores : [{ type : mongoose.Schema.ObjectId, ref : 'Store' }]  
})
```

//populate('store') will replace the store Id with the corresponding store object
Book.find({}).populate('stores')

Aggregation Queries





Aggregation Queries

- Summarize data typically for reports
- How would we solve this in SQL?

SELECT **GROUP BY** HAVING

- What About MongoDB?

=> **Aggregation Pipeline**

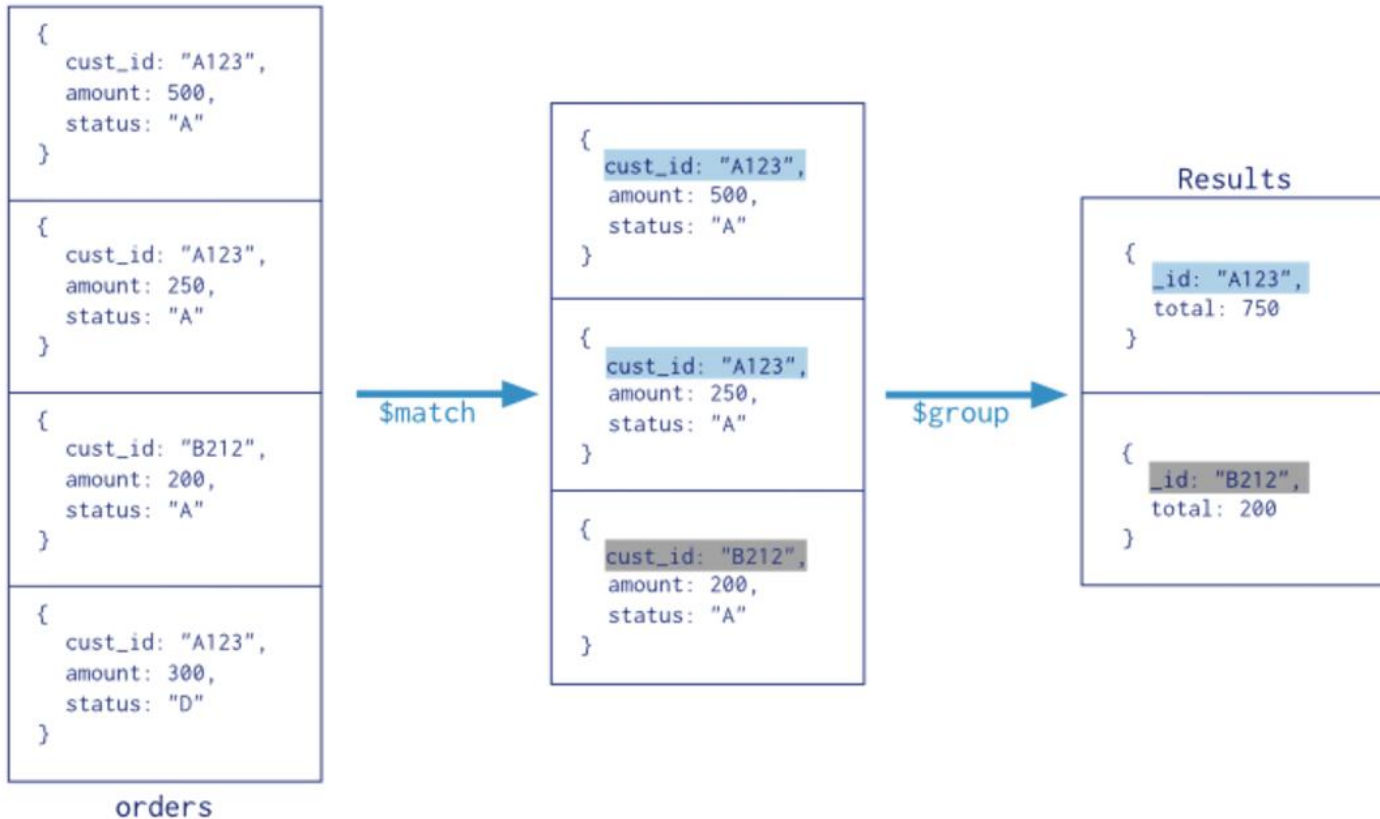


- Pipeline of functions to filter, group, and sort documents
- Operations executed in sequential order
- Output of one stage is used as an input of next

Aggregation Pipeline

- Allows filtering then grouping documents by specific fields

Collection
↓
`db.orders.aggregate([`
 `$match stage → { $match: { status: "A" } },`
 `$group stage → { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }`
 `]`)



Pipeline Operators

- `$match` Filter documents
- `$group` Summarize documents
- `$sort` Order documents
- `$limit` Limit returned results

- **`$group`** specifies:
 - Properties to group by
 - Computed output properties using `$max`, `$min`, `$avg`, `$sum` ...

\$group Examples

- Return average GPA for all students

```
Student.aggregate([
  {
    "$group" : { "_id" : null,
    "avgGPA" : {$avg : "$gpa"}
  }}])
```

- Return total completed Credit Hours per student

```
StudentCourse.aggregate([
  {
    "$group" : { "_id" : studentId,
    "completedCHs" : {$sum : "$CourseCH"}
  }}])
```

Resources

- Mongoose Documentation

<http://mongoosejs.com/docs/>

- Queries Cheat Sheet

http://s3.amazonaws.com/info-mongodb-com/mongodb_qrc_queries.pdf

- Aggregation Queries

<https://docs.mongodb.com/manual/reference/operator/aggregation/group/>